

Use `kubectl options` for a list of global command line options (applies to all commands).

• @Subash-1507 → /workspaces/Qwen-Chatbot (main) \$ kubectl get namespaces --all-namespaces

NAME	STATUS	AGE
chatbot	Active	10m
default	Active	17m
kube-node-lease	Active	17m
kube-public	Active	17m
kube-system	Active	17m

• @Subash-1507 → /workspaces/Qwen-Chatbot (main) \$ kubectl get pods --all-namespaces

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
chatbot	backend-856c758f4f-hktnj	1/1	Running	0	11m
chatbot	frontend-577fdd844b-bp5bz	1/1	Running	0	11m
kube-system	coredns-559f6c778d-vcz14	1/1	Running	0	17m
kube-system	etcd-minikube	1/1	Running	0	17m
kube-system	kindnet-2csbh	1/1	Running	0	17m
kube-system	kube-apiserver-minikube	1/1	Running	0	17m
kube-system	kube-controller-manager-minikube	1/1	Running	0	17m
kube-system	kube-proxy-bsg6j	1/1	Running	0	17m
kube-system	kube-scheduler-minikube	1/1	Running	0	17m
kube-system	metrics-server-59959bcbfd-cwx6v	1/1	Running	0	2m21s
kube-system	storage-provisioner	1/1	Running	0	17m

○ @Subash-1507 → /workspaces/Qwen-Chatbot (main) \$

pod "metrics-server-59959bcbfd-cwx6v" deleted from kube-system namespace

• @Subash-1507 → /workspaces/Qwen-Chatbot (main) \$ kubectl get pods,hpa,svc -n chatbot -o wide

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
pod/backend-856c758f4f-hktnj	1/1	Running	1 (3m11s ago)	41m	10.244.0.2	minikube	<none>	<none>
pod/frontend-577fdd844b-bp5bz	1/1	Running	1 (3m12s ago)	41m	10.244.0.4	minikube	<none>	<none>

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
horizontalpodautoscaler.autoscaling/backend-hpa	Deployment/backend	cpu: <unknown>/60%	1	3	1	41m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
service/backend	ClusterIP	10.96.247.215	<none>	5001/TCP	41m	app=backend
service/frontend	NodePort	10.104.39.43	<none>	80:30080/TCP	41m	app=frontend

• @Subash-1507 → /workspaces/Qwen-Chatbot (main) \$

backend-hpa Deployment/backend cpu: 0%/60% 1 3 1 42m

○ @Subash-1507 → /workspaces/Qwen-Chatbot (main) \$ kubectl port-forward --address 0.0.0.0 -n chatbot svc/frontend 8080:80

Forwarding from 0.0.0.0:8080 -> 5000

Handling connection for 8080

Handling connection for 8080

Qwen Chatbot

Powered by Qwen2.5-0.5B-Instruct · running offline in Kubernetes

hello

You

AI

Hello! How can I help you today? Is there something specific you would like to know or discuss? I'm here to assist with any questions you may have.

pod: backend-856c758f4f-hktrj

Type a message... (Enter to send, Shift+Enter for newline)

Send

storage

30 GB gp3

```
Problems Output Debug Console Terminal Ports 1 bash + - [ ] [ ] ... [ ] [ ] X
@Subash-1507 → /workspaces/Qwen-Chatbot (main) $ kubectl describe pod backend-856c758f4f-hktnj -n chatbot
PodReadyToStartContainers True
Initialized True
Ready True
ContainersReady True
PodScheduled True
Volumes:
  kube-api-access-c4kqs:
    Type: Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    Optional: false
    DownwardAPI: true
QoS Class: Burstable
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
              node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type      Reason             Age   From          Message
  ----      -
Warning    FailedScheduling   15m   default-scheduler 0/1 nodes are available: 1 node(s) had intolerated taint(s). pr
emption: 0/1 nodes are available: 1 Preemption is not helpful for scheduling.
Normal     Scheduled          11m   default-scheduler  Successfully assigned chatbot/backend-856c758f4f-hktnj to minik
ube
Normal     Pulled             11m   kubelet        spec.containers{backend}: Container image "chatbot-backend:late
st" already present on machine and can be accessed by the pod
Normal     Created            11m   kubelet        spec.containers{backend}: Container created
Normal     Started            11m   kubelet        spec.containers{backend}: Container started
Warning    Unhealthy          8m11s  kubelet        spec.containers{backend}: Readiness probe failed: Get "http://1
0.244.0.4:5001/health": context deadline exceeded (Client.Timeout exceeded while awaiting headers)
@Subash-1507 → /workspaces/Qwen-Chatbot (main) $
```

```
Problems Output Debug Console Terminal Ports 1 bash + - [ ] [ ] ... [ ] [ ] X
@Subash-1507 → /workspaces/Qwen-Chatbot (main) $ kubectl describe pod frontend-577fdd844b-bp5bz -n chatbot
Conditions:
  Type      Status
PodReadyToStartContainers True
Initialized True
Ready True
ContainersReady True
PodScheduled True
Volumes:
  kube-api-access-n5kdm:
    Type: Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    Optional: false
    DownwardAPI: true
QoS Class: Burstable
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
              node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type      Reason             Age   From          Message
  ----      -
Warning    FailedScheduling   16m   default-scheduler 0/1 nodes are available: 1 node(s) had intolerated taint(s). pr
emption: 0/1 nodes are available: 1 Preemption is not helpful for scheduling.
Normal     Scheduled          12m   default-scheduler  Successfully assigned chatbot/frontend-577fdd844b-bp5bz to minik
ube
Normal     Pulled             12m   kubelet        spec.containers{frontend}: Container image "chatbot-frontend:lat
est" already present on machine and can be accessed by the pod
Normal     Created            12m   kubelet        spec.containers{frontend}: Container created
Normal     Started            12m   kubelet        spec.containers{frontend}: Container started
@Subash-1507 → /workspaces/Qwen-Chatbot (main) $
```

```

D @Subash-1507 → /workspaces/Qwen-Chatbot (main) $ kubectl logs -n chatbot deployment/backend -f
[2026-09-26 11:04:54 +0000] [1] [INFO] Starting gunicorn 26.2.0
[2026-09-26 11:04:54 +0000] [1] [INFO] Listening at: http://0.0.0.0:5001 (1)
[2026-09-26 11:04:54 +0000] [1] [INFO] Using worker: sync
[2026-09-26 11:04:54 +0000] [7] [INFO] Booting worker with pid: 7
[2026-09-26 11:04:54 +0000] [1] [INFO] Control socket listening at /root/.gunicorn/gunicorn.ctl
[backend] loading Qwen/Qwen2.5-0.5B-Instruct ...
[transformers] `torch_dtype` is deprecated! Use `dtype` instead!
Loading weights: 100%|██████████| 290/290 [00:00<00:00, 4526.82it/s]
[backend] ready.
█

```